

Introduction

Dienstag, 1. November 2022 12:02

- Convolutions (standard vision approach) not applicable to asynchronous data
 - o Create synchronous approach by "leaking" information from past
 - Asynchronous (event-by-event) or Synchronous (batches) input, synchronous output (fixed rate of reconstruction)
 - Example: CNN accepting async event, processing events -> events/s = event throughput
 - Single-layer CNN about 1 million events/s (cameras have up to 10 million events/s)
 - Problem: event-rate above event throughput -> no real time possible
 - o Proposed method: sync/async hybrid
 - Problem with async event-by-event: 1 output per input event
 - Idea: apply only correction factor to some previous convolution result, performed as multi-thread process -> does not hinder event-throughput
 - Computing of correction factor (fast) and convolution update (slow) decoupled
 - Small cost in precision

Reference Implementation

Dienstag, 1. November 2022 12:29

- Each event requires internal representation C^* to be updated for kernelSize around event location $x_k \rightarrow$ patch wise updates, out-of-sync (denoted by *)
- 3 components per update:
 - Integration of event polarity
 - Temporal decay
 - user-defined scene parameter α
 - Time between current and prev. update C^* , stored in matrix T (last update time/pixel)
 - Convolution with kernel W
- Algorithm:
 1. *for* $(u, v) \in \text{patch}(x_i, y_i)$ *do*
 2. $dt \leftarrow t_i - T(u, v);$
 3. $C^*(u, v) \leftarrow C^*(u, v)e^{-\alpha dt} + W(u, v)p_i s;$
 4. $T(u, v) \leftarrow t_i;$
 5. $e_i \leftarrow \{x_i, y_i, t_i, C^*(x_i, y_i)\};$

High-throughput implementation

Dienstag, 1. November 2022 12:29

- Produces async output $\tilde{e}_i(x_i, y_i, t_i, c_i)$, reduces per-event computation by decoupling main operations
- Output event gets convolution score from convolution image from a secondary process
- Event-throughput maximised, cause secondary process multi-threaded
- c_i is approximation, cause convolved image production delayed
 - o Correction factor v added
- Event-throughput from $0.61 * 10^6 \frac{e}{s}$ to $9.77 * 10^6 \frac{e}{s}$ (significant increase)(using $kernelSize = 3$)
- Latency of the order of milliseconds on benchmark data set ($35 * 10^6 \frac{e}{s}$) for high-throughput (seconds for reference)
 - o Real-time possible, delay does not increase over time
- Average error from reference implementation $< 0.5\%$
- Github: <https://github.com/event-driven-robotics/high-throughput-convolutions>
- Video: doi